







B. Whole-client-deletion survivors | group quality across refinement rounds

One fixed client removed per dataset: 300 / 2,259 / 299 records; $k = 10$; five seeds (10000-10004).

Group 0 Group 1 Worst-group cost Population-average cost



Thin paths: all five paired seeds. Bold markers: mean across seeds. Bars: $\pm$ one sample SD ($n - 1$).

Worst-group curve = mean of each seed's maximum group cost; population curve = mean of each seed's count-weighted cost.

Axes show absolute cost with nonzero origins; dataset scales differ. Group masks are frozen; Adult/Credit raw provenance is incomplete.

Adult, T=1 to T=2: group 1 cost rises in 2/5 seeds while worst-group cost falls. No ordinary-Lloyd baseline is shown.